

Sandeep Pandit (TAS366)

✓ Customer Purchasing Behavior Analysis Using Spark

```
!pip install pyspark matplotlib seaborn pandas
```

```
Requirement already satisfied: pyspark in /usr/local/lib/python3.11/dist-packages (3.5.5)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.11/dist-packages (from pyspark) (0.10.9.7)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.
Requirement already satisfied: numpy>=1.23 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->ma
```

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, sum, count, desc, to_date, avg, when, month, year, hour
from pyspark.sql.types import StringType
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
spark = SparkSession.builder.appName("CustomerPurchasingAnalysis").getOrCreate()
```

```
spark.conf.set("spark.sql.legacy.timeParserPolicy", "LEGACY")
```

```
data = [("7/29/10 18:49"), ("12/5/15 08:30"), ("3/15/20 14:15")]
columns = ["timestamp_col"]
```

```
# Load datasets
file1_path = "/content/file1.csv" # Update with actual path
file2_path = "/content/file2.csv" # Update with actual path
```

```
df1 = spark.read.csv(file1_path, header=True, inferSchema=True)
df1.show()
df2 = spark.read.csv(file2_path, header=True, inferSchema=True)
df2.show()
```

```

+-----+-----+-----+-----+-----+-----+-----+
|Invoice|StockCode|Description|Quantity| InvoiceDate|Price|Customer ID|Country|
+-----+-----+-----+-----+-----+-----+-----+
| 489434| 85048|15CM CHRISTMAS GL...|12|12/1/09 7:45|6.95|13085|United Kingdom|
| 489434| 79323P| PINK CHERRY LIGHTS|12|12/1/09 7:45|6.75|13085|United Kingdom|
| 489434| 79323W| WHITE CHERRY LIGHTS|12|12/1/09 7:45|6.75|13085|United Kingdom|
| 489434| 22041|"RECORD FRAME 7""...|48|12/1/09 7:45|2.1|13085|United Kingdom|
| 489434| 21232|STRAWBERRY CERAMI...|24|12/1/09 7:45|1.25|13085|United Kingdom|
| 489434| 22064|PINK DOUGHNUT TRI...|24|12/1/09 7:45|1.65|13085|United Kingdom|
| 489434| 21871| SAVE THE PLANET MUG|24|12/1/09 7:45|1.25|13085|United Kingdom|
| 489434| 21523|FANCY FONT HOME S...|10|12/1/09 7:45|5.95|13085|United Kingdom|
| 489435| 22350| CAT BOWL |12|12/1/09 7:46|2.55|13085|United Kingdom|
| 489435| 22349|DOG BOWL , CHASIN...|12|12/1/09 7:46|3.75|13085|United Kingdom|
| 489435| 22195|HEART MEASURING S...|24|12/1/09 7:46|1.65|13085|United Kingdom|
| 489435| 22353|LUNCHBOX WITH CUT...|12|12/1/09 7:46|2.55|13085|United Kingdom|
| 489436| 48173C|DOOR MAT BLACK FL...|10|12/1/09 9:06|5.95|13078|United Kingdom|
| 489436| 21755|LOVE BUILDING BLO...|18|12/1/09 9:06|5.45|13078|United Kingdom|
| 489436| 21754|HOME BUILDING BLO...|3|12/1/09 9:06|5.95|13078|United Kingdom|
| 489436| 84879|ASSORTED COLOUR B...|16|12/1/09 9:06|1.69|13078|United Kingdom|
| 489436| 22119| PEACE WOODEN BLO...|3|12/1/09 9:06|6.95|13078|United Kingdom|
| 489436| 22142|CHRISTMAS CRAFT W...|12|12/1/09 9:06|1.45|13078|United Kingdom|
| 489436| 22296|HEART IVORY TRELL...|12|12/1/09 9:06|1.65|13078|United Kingdom|
| 489436| 22295|HEART FILIGREE DO...|12|12/1/09 9:06|1.65|13078|United Kingdom|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
536365	85123A	WHITE HANGING HEA...	6	12/1/10 8:26	2.55	17850	United Kingdom
536365	71053	WHITE METAL LANTERN	6	12/1/10 8:26	3.39	17850	United Kingdom
536365	84406B	CREAM CUPID HEART...	8	12/1/10 8:26	2.75	17850	United Kingdom
536365	84029G	KNITTED UNION FLA...	6	12/1/10 8:26	3.39	17850	United Kingdom
536365	84029E	RED WOOLLY HOTTIE...	6	12/1/10 8:26	3.39	17850	United Kingdom
536365	22752	SET 7 BABUSHKA NE...	2	12/1/10 8:26	7.65	17850	United Kingdom
536365	21730	GLASS STAR FROSTE...	6	12/1/10 8:26	4.25	17850	United Kingdom
536366	22633	HAND WARMER UNION...	6	12/1/10 8:28	1.85	17850	United Kingdom
536366	22632	HAND WARMER RED P...	6	12/1/10 8:28	1.85	17850	United Kingdom
536368	22960	JAM MAKING SET WI...	6	12/1/10 8:34	4.25	13047	United Kingdom
536368	22913	RED COAT RACK PAR...	3	12/1/10 8:34	4.95	13047	United Kingdom
536368	22912	YELLOW COAT RACK ...	3	12/1/10 8:34	4.95	13047	United Kingdom
536368	22914	BLUE COAT RACK PA...	3	12/1/10 8:34	4.95	13047	United Kingdom
536367	84879	ASSORTED COLOUR B...	32	12/1/10 8:34	1.69	13047	United Kingdom
536367	22745	POPPY'S PLAYHOUSE...	6	12/1/10 8:34	2.1	13047	United Kingdom
536367	22748	POPPY'S PLAYHOUSE...	6	12/1/10 8:34	2.1	13047	United Kingdom
536367	22749	FELTCRAFT PRINCES...	8	12/1/10 8:34	3.75	13047	United Kingdom
536367	22310	IVORY KNITTED MUG...	6	12/1/10 8:34	1.65	13047	United Kingdom
536367	84969	BOX OF 6 ASSORTED...	6	12/1/10 8:34	4.25	13047	United Kingdom
536367	22623	BOX OF VINTAGE JI...	3	12/1/10 8:34	4.95	13047	United Kingdom

only showing top 20 rows

```
df = df1.union(df2)
```

```
# Display schema and first few rows
```

```
df.printSchema()
```

```
df.show(5)
```

```
root
|-- Invoice: string (nullable = true)
|-- StockCode: string (nullable = true)
|-- Description: string (nullable = true)
|-- Quantity: integer (nullable = true)
|-- InvoiceDate: string (nullable = true)
|-- Price: double (nullable = true)
|-- Customer ID: integer (nullable = true)
|-- Country: string (nullable = true)
```

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
489434	85048	15CM CHRISTMAS GL...	12	12/1/09 7:45	6.95	13085	United Kingdom
489434	79323P	PINK CHERRY LIGHTS	12	12/1/09 7:45	6.75	13085	United Kingdom
489434	79323W	WHITE CHERRY LIGHTS	12	12/1/09 7:45	6.75	13085	United Kingdom
489434	22041	"RECORD FRAME 7""...	48	12/1/09 7:45	2.1	13085	United Kingdom
489434	21232	STRAWBERRY CERAMI...	24	12/1/09 7:45	1.25	13085	United Kingdom

only showing top 5 rows

```
# Data Cleaning: Removing negative or zero quantities
```

```
df_filtered = df.filter(col("Quantity") > 0)
```

```
# Fix column names
```

```
df_filtered = df_filtered.withColumnRenamed("Customer ID", "CustomerID")
```

```
# Convert CustomerID to StringType
```

```
df_filtered = df_filtered.withColumn("CustomerID", col("CustomerID").cast(StringType()))
```

```
# Mapping CustomerID to Total Spend
```

```
df_spending = df_filtered.groupBy("CustomerID").agg(
    sum(col("Quantity") * col("Price")).alias("TotalSpend")
)
```

```
# Case 1: Excluding NULLs from analysis
```


```
df_spending_no_nulls = df_spending.filter(col("CustomerID").isNotNull())
```

```
# Case 2: Including "Unknown" instead of NULL
```

```
df_filtered = df_filtered.fillna({"CustomerID": "Unknown"})
```

```
df_spending_with_unknown = df_spending.fillna({"CustomerID": "Unknown"})
```

```
# Identifying top 10 customers by total spending
top_customers_no_nulls = df_spending_no_nulls.orderBy(desc("TotalSpend")).limit(10)
top_customers_no_nulls.show()
top_customers=top_customers_no_nulls
top_customers_with_unknown = df_spending_with_unknown.orderBy(desc("TotalSpend")).limit(10)
top_customers_with_unknown.show()
```




CustomerID	TotalSpend
18102	608821.6500000001
14646	528602.5199999989
14156	313946.3699999996
14911	295972.62999999954
17450	246973.08999999997
13694	196482.8100000001
17511	175603.54999999984
16446	168472.5
16684	147142.77
12415	144458.3699999998

CustomerID	TotalSpend	Segment
Unknown	3070862.8200000157	High Value Customers
18102	608821.6500000001	High Value Customers
14646	528602.5199999989	High Value Customers
14156	313946.3699999996	High Value Customers
14911	295972.62999999954	High Value Customers
17450	246973.08999999997	High Value Customers
13694	196482.8100000001	High Value Customers
17511	175603.54999999984	High Value Customers
16446	168472.5	High Value Customers
16684	147142.77	High Value Customers


```
# Convert to Pandas for visualization
top_customers_no_nulls_pd = top_customers_no_nulls.toPandas()
top_customers_with_unknown_pd = top_customers_with_unknown.toPandas()
```

```
# Visualization: Top 10 Customers (Excluding NULLs)
plt.figure(figsize=(10, 5))
sns.barplot(x="CustomerID", y="TotalSpend", data=top_customers_no_nulls_pd, palette="coolwarm")
plt.xticks(rotation=45)
plt.xlabel("Customer ID")
plt.ylabel("Total Spending")
plt.title("Top 10 Customers by Total Spending (Excluding NULLs)")
plt.show()
```

 <ipython-input-365-f07a143c930e>:3: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to

```
sns.barplot(x="CustomerID", y="TotalSpend", data=top_customers_no_nulls_pd, palette="coolwarm")
```



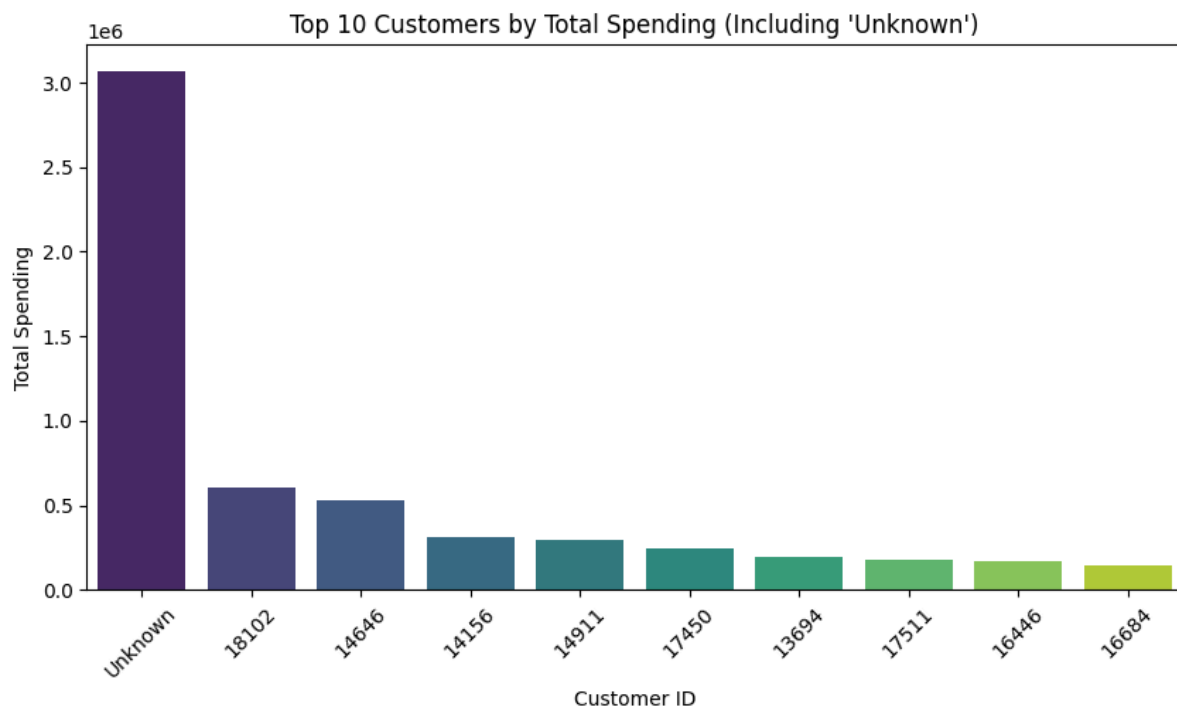
Visualization: Top 10 Customers (Including 'Unknown' for NULLs)

```
plt.figure(figsize=(10, 5))
sns.barplot(x="CustomerID", y="TotalSpend", data=top_customers_with_unknown_pd, palette="viridis")
plt.xticks(rotation=45)
plt.xlabel("Customer ID")
plt.ylabel("Total Spending")
plt.title("Top 10 Customers by Total Spending (Including 'Unknown')")
plt.show()
```

 <ipython-input-366-f369a3b1ffbd>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to


```
sns.barplot(x="CustomerID", y="TotalSpend", data=top_customers_with_unknown_pd, palette="viridis")
```



```
# Determining the most frequently purchased products
top_products = df_filtered.groupBy("StockCode", "Description").agg(
    count("Quantity").alias("PurchaseCount")
).orderBy(desc("PurchaseCount")).limit(10)
```

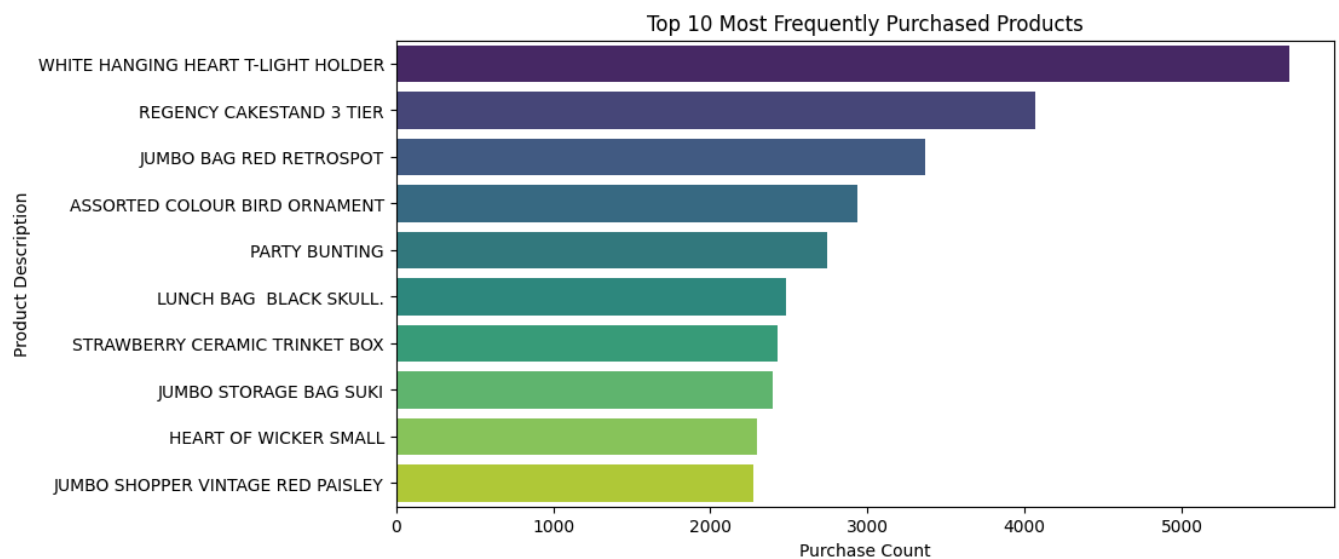
```
# Convert to Pandas for visualization
top_products_pd = top_products.toPandas()
```

```
# Visualization: Most Frequently Purchased Products
plt.figure(figsize=(10, 5))
sns.barplot(x="PurchaseCount", y="Description", data=top_products_pd, palette="viridis")
plt.xlabel("Purchase Count")
plt.ylabel("Product Description")
plt.title("Top 10 Most Frequently Purchased Products")
plt.show()
```

 <ipython-input-368-4808a7e78d5e>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to

```
sns.barplot(x="PurchaseCount", y="Description", data=top_products_pd, palette="viridis")
```




```
# Analyzing purchasing trends over time
df_filtered = df_filtered.withColumn("InvoiceDate", to_date(col("InvoiceDate"), "MM/dd/yy HH:mm"))
```

```
# Extract year and month
df_filtered = df_filtered.withColumn("Year", year(col("InvoiceDate")))
df_filtered = df_filtered.withColumn("Month", month(col("InvoiceDate")))
```

```
monthly_trends = df_filtered.groupBy("Year", "Month").agg(
    sum("Quantity").alias("TotalSales")
).orderBy("Year", "Month")
monthly_trends.show()
monthly_trends_pd = monthly_trends.toPandas()
```

```
# Create a 'Year-Month' column for better visualization
monthly_trends_pd["Year-Month"] = monthly_trends_pd["Year"].astype(str) + "-" + monthly_trends_pd["Month"].astype(str)
```

 +-----+
|Year|Month|TotalSales|
+-----+
2009	12	445861
2010	1	396087
2010	2	392763
2010	3	531689
2010	4	387241
2010	5	425348
2010	6	415541
2010	7	359616
2010	8	522803
2010	9	594108
2010	10	624378
2010	11	733705
2010	12	544764


2011	1	397716
2011	2	286695
2011	3	384950
2011	4	312176
2011	5	399425
2011	6	394337
2011	7	407539

+-----+

only showing top 20 rows

```
# Visualization: Monthly Sales Trends for Both Years
plt.figure(figsize=(12, 6))
sns.lineplot(data=monthly_trends_pd, x="Year-Month", y="TotalSales", hue="Year", marker="o", palette="Set1")

plt.xticks(rotation=45)
plt.xlabel("Month-Year")
plt.ylabel("Total Sales")
plt.title("Monthly Sales Trends (2009-2010 and 2010-2011)")
plt.grid()
plt.legend(title="Year")
plt.show()
```



```
# Customer Segmentation based on Spending Behavior
df_spending_with_unknown = df_spending_with_unknown.withColumn(
    "Segment",
    when(col("TotalSpend") > 5000, "High Value Customers")
    .when(col("TotalSpend") > 1000, "Mid Value Customers")
    .otherwise("Low Value Customers")
)


customer_segments_pd = df_spending_with_unknown.groupBy("Segment").count().toPandas()

# Visualization: Customer Segmentation
plt.figure(figsize=(8, 8))
plt.pie(
    customer_segments_pd["count"],
    labels=customer_segments_pd["Segment"],
    autopct="%1.1f%%",
    colors=sns.color_palette("Set2"),
    startangle=140,
    wedgeprops={"edgecolor": "black"}
)
```

```
plt.title("Customer Segmentation Based on Spending Behavior")
plt.show()
```



Customer Segmentation Based on Spending Behavior



```
# Customer Retention Analysis
customer_orders = df_filtered.groupBy("CustomerID").agg(count("InvoiceDate").alias("OrderCount"))
repeat_customers = customer_orders.filter(col("OrderCount") > 1).count()
one_time_customers = customer_orders.filter(col("OrderCount") == 1).count()

total_customers = repeat_customers + one_time_customers
retention_rate = repeat_customers / total_customers
print(f"Customer Retention Rate: {retention_rate * 100:.2f}%")
```



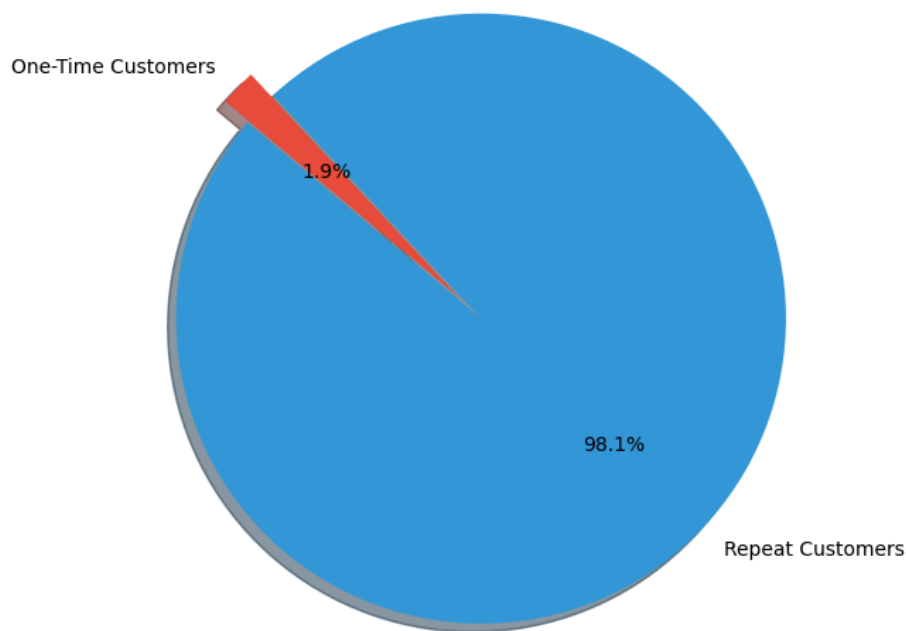
Customer Retention Rate: 98.06%

```
# Plot the Retention Rate
labels = ['Repeat Customers', 'One-Time Customers']
sizes = [repeat_customers, one_time_customers]
colors = ['#3498db', '#e74c3c']
explode = (0.1, 0) # Explode the first slice for emphasis

plt.figure(figsize=(7, 7))
plt.pie(sizes, labels=labels, autopct='%1.1f%%', colors=colors, explode=explode, startangle=140, shadow=True)
plt.title("Customer Retention Analysis")
plt.show()
```



Customer Retention Analysis



```
# Extract month from InvoiceDate
df_filtered = df_filtered.withColumn("Month", month(col("InvoiceDate")))


# Count unique customers per month
monthly_customers = df_filtered.groupBy("Month").agg(count("CustomerID").alias("UniqueCustomers")).orderBy("Month")

# Convert to Pandas for visualization
monthly_customers_pd = monthly_customers.toPandas()

# Line chart of monthly customer retention trends
plt.figure(figsize=(10, 5))
sns.lineplot(data=monthly_customers_pd, x="Month", y="UniqueCustomers", marker="o", color="b")
plt.xlabel("Month")
plt.ylabel("Number of Unique Customers")
plt.title("Monthly Customer Retention Trends")
plt.grid()
plt.show()
```



Monthly Customer Retention Trends




```

# Extract Year from InvoiceDate
df_filtered = df_filtered.withColumn("Year", year(col("InvoiceDate")))

# Filter data for two time periods
df_2009_10 = df_filtered.filter((col("Year") == 2009) | (col("Year") == 2010))
df_2010_11 = df_filtered.filter((col("Year") == 2010) | (col("Year") == 2011))

# Count unique customers
customers_2009_10 = df_2009_10.select("CustomerID").distinct().count()
customers_2010_11 = df_2010_11.select("CustomerID").distinct().count()


# Find retained customers (customers in both periods)
common_customers = df_2009_10.select("CustomerID").distinct().intersect(
    df_2010_11.select("CustomerID").distinct()
).count()

# Retention rate calculation
retention_2009_10 = common_customers / customers_2009_10 * 100
retention_2010_11 = common_customers / customers_2010_11 * 100

# Plotting retention rate
import matplotlib.pyplot as plt
import seaborn as sns


plt.figure(figsize=(7, 5))
sns.barplot(x=["2009-10", "2010-11"], y=[retention_2009_10, retention_2010_11], palette="Blues")
plt.ylabel("Retention Rate (%)")
plt.title("Customer Retention Rate Comparison")
plt.ylim(0, 100)
plt.show()

```

 <ipython-input-378-96d50c701547>:19: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` to silence this warning.

```
sns.barplot(x=["2009-10", "2010-11"], y=[retention_2009_10, retention_2010_11], palette="Blues")
```



```

from pyspark.sql.functions import countDistinct, col, month

# Get monthly unique customers for each period
df_2009_10 = df_2009_10.withColumn("Month", month(col("InvoiceDate")))
df_2010_11 = df_2010_11.withColumn("Month", month(col("InvoiceDate")))

# Group by Month and count distinct customers
monthly_customers_2009_10 = df_2009_10.groupBy("Month").agg(countDistinct("CustomerID").alias("UniqueCustomers"))
monthly_customers_2010_11 = df_2010_11.groupBy("Month").agg(countDistinct("CustomerID").alias("UniqueCustomers"))


# Convert to Pandas for plotting
monthly_customers_2009_10_pd = monthly_customers_2009_10.toPandas()
monthly_customers_2010_11_pd = monthly_customers_2010_11.toPandas()

# Plot line chart

```

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10, 5))
sns.lineplot(data=monthly_customers_2009_10_pd, x="Month", y="UniqueCustomers", marker="o", label="2009-10", color="red")
sns.lineplot(data=monthly_customers_2010_11_pd, x="Month", y="UniqueCustomers", marker="o", label="2010-11", color="blue")
plt.xlabel("Month")
plt.ylabel("Number of Unique Customers")
plt.title("Monthly Customer Activity Comparison")
plt.legend()
plt.grid()
plt.show()
```



```
from pyspark.sql.functions import sum


# Calculate total sales for both years
sales_2009_10 = df_2009_10.agg(sum(col("Quantity") * col("Price"))).collect()[0][0]
sales_2010_11 = df_2010_11.agg(sum(col("Quantity") * col("Price"))).collect()[0][0]

# Plot bar chart for total sales comparison
plt.figure(figsize=(7, 5))
sns.barplot(x=["2009-10", "2010-11"], y=[sales_2009_10, sales_2010_11], palette="Greens")
plt.ylabel("Total Sales ($)")
plt.title("Total Sales Comparison")
plt.show()
```

 <ipython-input-380-2aae60ff4b63>:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to

```
sns.barplot(x=["2009-10", "2010-11"], y=[sales_2009_10, sales_2010_11], palette="Greens")
```



```
# Aggregate revenue per customer
customer_revenue = df_filtered.groupby("CustomerID").agg(sum(col("Quantity") * col("Price")).alias("TotalRevenue"))

# Segment customers into groups
customer_revenue = customer_revenue.withColumn(
    "Segment",
    when(col("TotalRevenue") < 500, "Low Spender")
    .when((col("TotalRevenue") >= 500) & (col("TotalRevenue") < 2000), "Medium Spender")
    .otherwise("High Spender")
)

# Count customers in each segment
segment_distribution = customer_revenue.groupBy("Segment").count().toPandas()

# Pie chart visualization
plt.figure(figsize=(8, 8))
plt.pie(
    segment_distribution["count"],
    labels=segment_distribution["Segment"],
    autopct="%1.1f%%",
    colors=sns.color_palette("Set2"),
    startangle=140,
    wedgeprops={"edgecolor": "black"}
)
plt.title("Revenue Contribution by Customer Segment")
plt.show()
```



Revenue Contribution by Customer Segment

